

# Kubernetes Lab Notes: Networking & Persistent Storage

## Table of Contents

1. Lab Setup
  2. Pod & Node Networking
  3. Service Discovery
  4. Kubernetes Storage
  5. StatefulSets
- 

## 1. Lab Setup

### What We Have

- **CCE Cluster** = A whole Kubernetes city
- **Node** = An apartment building in the city
- **ECS** = Your house outside the city

### Installing kubectl - The City Remote Control

**Q) Why install kubectl in ECS (your house) and not on the Node (apartment building)?**

- It's safer to control the city from your house than from inside an apartment building
- More convenient - you can manage everything from one place

## How to Get Your Remote Control Working:

1. Install `kubectl` on your computer
2. Get the special config file (like pairing your remote)
3. Put it in the right place: `~/.kube/config`
4. Test it: `kubectl get nodes` - this should show your buildings

```
# This is like testing your remote control
kubectl get nodes
```

## Sample Output:

NAME	STATUS	ROLES	AGE	VERSION
cce01-node1	Ready	<none>	15d	v1.25.0
cce01-node2	Ready	<none>	15d	v1.25.0

## 2. Pod & Node Networking

### What's a Pod?

A **Pod** is like a single shop in our city. Each shop has:

- Its own phone number (IP address)
- Can be temporary (opens and closes)

### Q) What do you mean by ephemeral?

- “Ephemeral” means **temporary** - like a pop-up shop
- Pods can close and reopen with **new phone numbers**
- This is why we need Services (phone books) later!

## Let's Test Communication

### Step 1: Open an Nginx Shop

```
kubectl run nginx --image=nginx --port=80
```

### Step 2: Find Its Phone Number

```
kubectl get pods -o wide
```

#### Sample Output:

NAME	READY	STATUS	RESTARTS	AGE	IP
NODE					
nginx	1/1	Running	0	10s	10.244.1.2
cce01-node1					

Phone number = 10.244.1.2

### Step 3: Call the Shop

```
curl 10.244.1.2
```

Q) When we write `curl <IP>`, what does it check?

- `curl` is like making a phone call 
- If someone answers (server works), you get the website
- If no one answers, you get an error

## Q) What was the purpose of doing all this?

- To prove that every shop (Pod) gets its own phone number (IP)
- And that everyone in the city can call each other directly!

## Testing Between Shops

What's Busybox? 

### Q) What is a busybox?

- Busybox is like a **Swiss Army knife** - lots of small useful tools
- Perfect for testing because it's simple and has everything we need

### Create a Busybox Shop and Call Nginx:

```
# Enter the busybox shop
kubectl run -it busybox --image=busybox -- /bin/sh

# Inside busybox, call nginx
telnet 10.244.1.2 80
```

### Q) What did we achieve by doing this?

- Proved that shops can call each other directly!
- No need for special phone operators

### 3. Service Discovery

#### The Big Problem

Shops (Pods) keep changing their phone numbers! How do you call “Bob’s Pizza” if their number changes daily?

#### The Solution: Services!

A **Service** is like a **city phone book** - it gives permanent numbers to groups of shops.

#### Creating Multiple Shops (Deployment)

`nginx-deploy.yaml`:

```
apiVersion: apps/v1 # This tells Kubernetes what
"language" we're speaking
kind: Deployment      # We're creating a shop
manager
metadata:
  name: nginx        # The manager's name
spec:
  replicas: 3        # Create 3 identical shops
  selector:          # The manager looks for shops
with this tag
    matchLabels:
      app: nginx
  template:           # This is the shop blueprint
    metadata:
      labels:
        app: nginx    # All shops get this tag
  spec:
    containers:
      - name: nginx
        image: nginx
      ports:
        - containerPort: 80
```

## Q) What does `selector` mean?

- It's how the manager finds "his" shops
- "Find all shops with tag `app=nginx`"

## Create the Shops:

```
kubectl apply -f nginx-deploy.yaml
kubectl get deployment
```

## Creating the Phone Book (Service)

### nginx-service.yaml:

```

apiVersion: v1
kind: Service      # This is a phone book
metadata:
  name: nginx-svc    # Phone book name
spec:
  selector:
    app: nginx      # List all shops with this tag
  ports:
    - protocol: TCP
      port: 8080      # Phone book number
      targetPort: 80  # Actual shop number

```

### Q) What does `selector: app: nginx` mean in the Service?

- “Put all shops with tag `app=nginx` in this phone book”

### Create the Phone Book:

```

kubectl apply -f nginx-service.yaml
kubectl get service

```

### Sample Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
PORT(S)	AGE		
nginx-svc	ClusterIP	10.105.32.105	<none>
8080/TCP	5m		

Permanent phone number = 10.105.32.105:8080

## Check Who's in the Phone Book:

```
kubectl get endpoints
```

### Sample Output:

```
nginx-svc    10.244.1.2:80,10.244.1.3:80,10.244.2.5:80
```

This shows the 3 actual shops behind the phone number!

## Making Your Shop Public

**Regular Service** = Internal company phone (only inside office)

**NodePort Service** = Public customer service line (anyone can call)

### nginx-expose.yaml:

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  type: NodePort      # Make it PUBLIC!
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
      nodePort: 30144 # Public number
```

## Q) What's the difference between normal service and expose.yaml?

- **Normal:** Only people in the city can call
- **NodePort:** Anyone in the world can call!

### Test the Public Number:

```
kubectl apply -f nginx-expose.yaml  
curl <ANY-NODE-IP>:30144
```

## Testing from Another Pod

### Q) What do you mean by “using a service to access workload from another pod”?

Let's create a customer and have them order pizza!

### Step 1: Create a Customer

```
# client.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: clientpod  
spec:  
  containers:  
    - name: clientpod  
      image: busybox:1.28.3  
      args:  
        - /bin/sh  
        - -c  
        - sleep 30000 # Stay alive for testing
```

## Step 2: Go Inside the Customer's Head

```
kubectl apply -f client.yaml  
kubectl exec -it clientpod -- /bin/sh
```

## Step 3: Look Up Pizza Shop in Phone Book

```
nslookup nginx-svc
```

This finds the permanent phone number

## Step 4: Call the Pizza Shop

```
wget -q -O - nginx-svc:8080
```

This successfully orders pizza! 🍕

**The Magic:** Even if actual pizza shops change, the phone book number stays the same!

---

## 4. Kubernetes Storage

**emptyDir = Shopping Cart **

**What it is:**

- Temporary storage while shopping
- **Gets emptied when you leave** (pod dies)

**empty-pod.yaml:**

```

apiVersion: v1
kind: Pod
metadata:
  name: em
spec:
  containers:
    - image: ubuntu
      name: test-container
      volumeMounts:
        - mountPath: /cache # Where to put cart in
          shop
          name: cache-volume
      args:
        - /bin/sh
        - -c
        - sleep 30000
    volumes:
      - name: cache-volume
        emptyDir: {} # This is the shopping cart

```

## Test it:

```

kubectl apply -f empty-pod.yaml
kubectl exec -it em -- /bin/bash

# Inside the container:
echo "Hello from lab!" > /cache/hello.file

```

## Q) Why `cat > hello.file<<EOF` ?

- It's a quick way to create a file with content
- Like writing a note and putting it in your cart

## hostPath = Writing on the Wall

### What it is:

- Store data directly on the building
- Stays even if your shop closes
- **Problem:** If shop moves to new building, writing isn't there!

### hostPath-pod.yaml:

```

volumes:
- name: hp-volume
  hostPath:
    path: /testdir      # Write on this wall
    type: Directory
  
```

## PV & PVC = Renting Storage Units

- **PV (PersistentVolume)** = Actual storage units in city
- **PVC (PersistentVolumeClaim)** = "I want to rent a storage unit!"

### Check Available Units:

```

kubectl get pv      # Show available storage units
kubectl get pvc     # Show rental requests
  
```

### Q) What's the difference between `kubectl get pv` and `kubectl get pvc` ?

- `get pv` = See all storage units in city
- `get pvc` = See who wants to rent units

### Using a Storage Unit in Your Shop:

**volumes:**

- `name`: pvc-volume
- `persistentVolumeClaim`:
  - `claimName`: cce-efs-import-mypvc # Your storage unit

**Best part:** Data survives even if shop closes and reopens!

---

# StatefulSets Explained Like Never Before!

The Problem: Why Food Trucks Don't Work for Banks 

Imagine you have **3 food trucks** (Regular Pods):

- Taco Truck #1
- Taco Truck #2
- Taco Truck #3

## What's the problem?

- They're all **identical** - same tacos, same menu
- If Taco Truck #1 breaks down, a **new identical one** replaces it
- Customers don't care which truck they get tacos from
- **Perfect for tacos, TERRIBLE for banks!**

# What ARE StatefulSets?

**StatefulSets are for businesses that NEED to remember things:**

Think of **3 bank branches**:

-  Bank Branch #1 (Downtown)
-  Bank Branch #2 (Uptown)
-  Bank Branch #3 (Suburbs)

**Each branch is SPECIAL:**

-  **Unique name** that NEVER changes
-  **Own vault** (storage) that NEVER mixes with others
-  **Opens/closes in specific order** (very organized!)

**Real examples:**

- Databases (MySQL, PostgreSQL)
  - Where each instance has **DIFFERENT data**
  - Applications that need to **remember who they are**
- 

Headless Service = Direct Employee Phone Lines 

**Regular Service (Main Company Number)**

 Call: 1-800-COMPANY

 You get: ANY available customer service rep

 You don't care WHO answers

## Headless Service (Direct Extensions)

- 📞 Call: Bob's direct line (x101)
- 👋 You get: BOB specifically
- 📞 Call: Alice's direct line (x102)
- 👋 You get: ALICE specifically

### Why do banks need this?

- You need to talk to **Bank Manager Bob** specifically
- You need to talk to **Loan Officer Alice** specifically
- Each has **different responsibilities and information**

### The Magic File:

```
apiVersion: v1
kind: Service
metadata:
  name: headless
spec:
  clusterIP: None # ⭐ THIS makes it HEADLESS!
  selector:
    app: tomcat
```

---

## Creating Our Organized Bank System

### Our Bank Blueprint:

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: tomcat-statefulset
spec:
  serviceName: tomcat-headless # Use direct phone
  lines
  replicas: 3 # Create 3 bank
  branches

  volumeClaimTemplates: # ★ Each branch
  gets OWN vault!
  - metadata:
      name: www-storage
    spec:
      resources:
        requests:
          storage: 5Gi
```

## What Gets Created:

tomcat-statefulset-0 www-storage-tomcat-statefulset-0

tomcat-statefulset-0.tomcat-headless...

tomcat-statefulset-1 www-storage-tomcat-statefulset-1

tomcat-statefulset-1.tomcat-headless...

tomcat-statefulset-2 www-storage-tomcat-statefulset-2

tomcat-statefulset-2.tomcat-headless...

**Key Point:** Each branch has its **OWN storage** that **never gets mixed up** with others!

# The MAGIC of StatefulSets ✨

## What if Disaster Strikes? 🔥

Bank Branch #1 burns down!

```
kubectl delete pod tomcat-statefulset-0  
kubectl get pod
```

## The MAGIC Happens:

1. **Same Name:** New branch opens as `tomcat-statefulset-0`  
(NOT a random name!)
2. **Same Vault:** Gets back the EXACT same vault:  
`www-storage-tomcat-statefulset-0`
3. **All Money Safe:** 💰 All customer data and money is STILL THERE!

This is IMPOSSIBLE with regular food trucks!

## Food Truck vs Bank Branch:

Truck #1 breaks New truck with RANDOM name

New branch with SAME name

Storage Loses all ingredients Keeps ALL customer money

Identity Just another taco truck Still Bank Branch #1

# Testing Our Direct Phone System

## Step 1: Create a Phone Tester

```
# Create a special phone book tester
kubectl run dnsutils --image=tutum/dnsutils --
command -- sleep infinity
```

## Step 2: Test the Direct Lines

```
# Look up all the direct phone numbers
kubectl exec dnsutils -- nslookup headless
```

### What you'll see:

```
tomcat-statefulset-0.tomcat-headless... → 10.244.1.10
tomcat-statefulset-1.tomcat-headless... → 10.244.1.11
tomcat-statefulset-2.tomcat-headless... → 10.244.1.12
```

**Now you can call each branch DIRECTLY!**

---

# Real-World Examples

## When to Use StatefulSets:

**Database Cluster** Each database has different data

**Message Queue** Each node has different messages

**Cache Cluster** Each instance has different cached data

**File Storage** Each pod manages different files

## When to Use Regular Deployments:

**Web Servers** All serve same website

**API Servers** All handle same requests

**Microservices** All provide same service

---

## Simple Summary

### Food Trucks (Regular Pods)

- All identical
- Interchangeable
- Lose everything if replaced

- **Good for:** Websites, APIs, stateless apps

## Bank Branches (StatefulSets)

- **Each is unique**
- **Keep their identity**
- **Preserve their data**
- **Good for:** Databases, queues, stateful apps

### The Golden Rule:

Use StatefulSets when your application needs to  
**remember who it is and what data it has!**

## Quick Test

**Question:** Should you use StatefulSet for a MySQL database cluster?

**Answer:**  **YES!** Each MySQL instance has different data and needs to keep its identity!

**Question:** Should you use StatefulSet for a website serving static pages?

**Answer:**  **NO!** All web servers are identical and interchangeable!

The beauty of StatefulSets is that they handle all the complexity of managing stateful applications while keeping everything organized and reliable! 

# Quick Summary Cheat Sheet

**Pod** Temporary shop Pop-up store

**Service** Permanent phone book Company phone number

**NodePort** Public phone number Customer service line

**emptyDir** Shopping cart Temporary storage

**PV/PVC** Storage unit rental Long-term storage

**Deployment** Interchangeable shops Food trucks

**StatefulSet** Organized branches Bank branches

**Headless Service** Direct extensions Employee direct lines

## Key Takeaways

1. **Pods are temporary** - they come and go with new IPs
2. **Services are permanent** - they provide stable access points
3. **Storage choices matter** - temporary vs permanent vs shared
4. **StatefulSets keep identity** - crucial for databases
5. **Kubernetes manages complexity** - you just describe what you want!

The beauty of Kubernetes is that it handles all the complicated stuff behind the scenes. You just describe your city, and Kubernetes builds and manages it for you! 🎉

**Remember:** Whether you're running a simple website or a complex database system, Kubernetes has the right tools for the job. Start simple, understand the basics, and gradually work your way up to more advanced concepts! 🚀